



NLP Sentiment Analysis

Ching and Chloe

TABLE OF CONTENTS



01

Business Question

04

**Models Creation
& Evaluation**

02

Data Collection

05

**Conclusion , Future
Improvements & Next
Steps**

03

Preprocessing



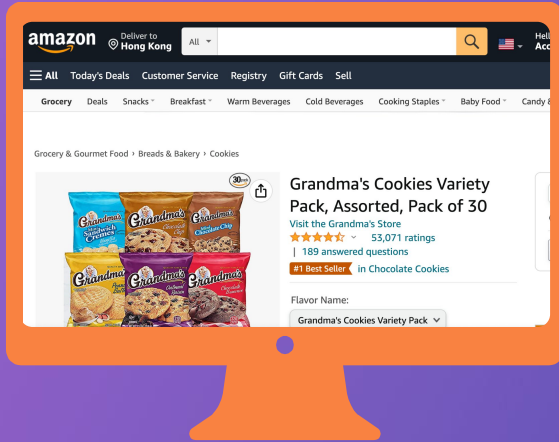


01

Business Question



Business Objectives



Boost Sales by better marketing tactics

By analysing customers' sentiment, company can know the trend and feature particular products

Improve company's products & services

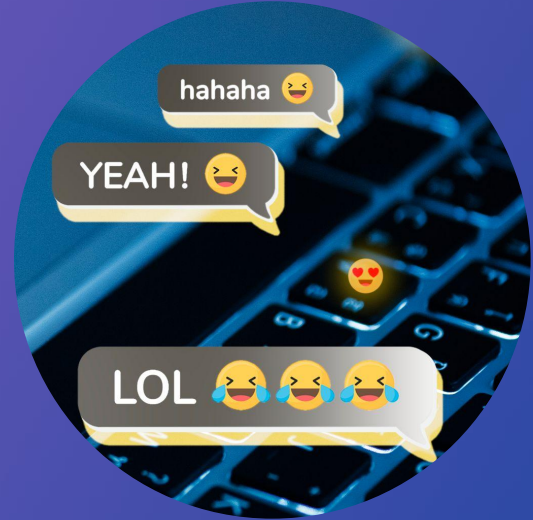
Customers may reviews some issues within the products or certain problems when using the services. Company could address those bugs and provide better one.

Business Objectives

Track sentiments in real-time

Retaining customers could generate more than half of the total revenue.

Hence by tracking real-time, company could quickly identify if sudden changes happened in customers' feeling. They could react promptly to retain the existing customers.



Maintain brand preception


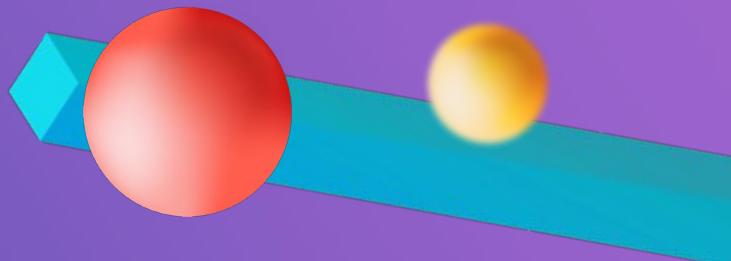
Customers value brand image. Yet, it requires long time to build the reputation. But at the same time,, it could be ruined within a single day. It is important to monitor the customers' feeling towards the products.



02

Data Collection

Data Collection



STANFORD NETWORK ANALYSIS
PROJECT · UPDATED 5 YEARS AGO


▲ 1876

New Notebook

Download (254 MiB)

Amazon Fine Food Reviews

Analyze ~500,000 food reviews from Amazon



<https://www.kaggle.com/datasets/snap/amazon-fine-food-reviews>



03

Data Preprocessing



Data Preprocessing

Original Dataset

```
df.head()
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	Text
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	5	1303862400	Good Quality Dog Food	I have bought several of the Vitality canned d...
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	1	1346976000	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	4	1219017600	"Delight" says it all	This is a confection that has been around a fe...
3	4	B000UA0QIQ	A395BORC6FGVXV	Karl	3	3	2	1307923200	Cough Medicine	If you are looking for the secret ingredient i...
4	5	B006K2ZZ7K	A1UQRSCLF8GW1T	Michael D. Bigham "M. Wassir"	0	0	5	1350777600	Great taffy	Great taffy at a great price. There was a wid...

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 568454 entries, 0 to 568453
```

```
Data columns (total 10 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	568454 non-null	int64
1	ProductId	568454 non-null	object
2	UserId	568454 non-null	object
3	ProfileName	568438 non-null	object
4	HelpfulnessNumerator	568454 non-null	int64
5	HelpfulnessDenominator	568454 non-null	int64
6	Score	568454 non-null	int64
7	Time	568454 non-null	int64
8	Summary	568427 non-null	object
9	Text	568454 non-null	object

```
dtypes: int64(5), object(5)
```

```
memory usage: 43.4+ MB
```

Data Preprocessing

1. Create a binary column and define score > 3 is positive, O.W. negative.

```
df['Positive'] = df['Score'].apply(lambda x: 1 if x > 3 else 0)
```

2. Drop duplicate rows by checking "UserID", "ProfileName", "Time" and "Text".

```
clean_df = df.drop_duplicates(subset={"UserID", "ProfileName", "Time", "Text"}, keep='first')
clean_df.shape

(393933, 11)
```

3. Only extract two columns (the comment itself, positive/negative)

```
df_text = clean_df[['Positive', 'Text']]
```

Text Preprocessing

1. Clean contractions

```
def clean(text):  
    text = re.sub(r"\n\t", " not", text)  
    text = re.sub(r"\re", " are", text)  
    text = re.sub(r"\s", " is", text)  
    text = re.sub(r"\d", " would", text)  
    text = re.sub(r"\ll", " will", text)  
    text = re.sub(r"\t", " not", text)  
    text = re.sub(r"\ve", " have", text)  
    text = re.sub(r"\m", " am", text)  
    return text
```

2. Define customized stopwords list

```
#customized stopwords list  
stop_words = set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've",  
"you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself',  
'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their',  
'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these', 'those',  
'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having', 'do', 'does',  
'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', 'while', 'of',  
'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during', 'before', 'after',  
'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under', 'again', 'further',  
'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'each', 'few', 'more',  
'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very',  
's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', 'm', 'o', 're',  
've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn',  
"hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't", 'mustn',  
"mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn', "wasn't", 'weren', "weren't",  
'won', "won't", 'wouldn', "wouldn't", "s", "..."])
```

3. Import SnowballStemmer

```
from nltk.stem.snowball import SnowballStemmer  
snow = SnowballStemmer('english')
```

Text Preprocessing

1. Use regex to clean:
 - (i) website links; (ii) words with numbers (e.g. 200ounces, 18yo); (iii) words with repeated letters (e.g. yummmmmmyyyy, ahhhhh)
2. Convert the texts into lower case and tokenize into words
3. Exclude stopwords
4. Stem the word tokens by Snowball Stemmer

```
def preprocessing(string):  
    word_sent = re.sub(r"http\S+", "", string)  
    word_sent = clean(word_sent)  
    word_sent = re.sub("\S*\d\S*", "", word_sent).strip() #removing words with numerical digits  
    word_sent = re.sub('[^A-Za-z]+', ' ', word_sent) # removing non-word characters  
    word_sent = re.sub(r'(\w)\1{2,}', r'\1', word_sent) #removing words of repeated letters  
    word_sent = word_tokenize(word_sent.lower())  
    word_sent = [word for word in word_sent if word not in stop_words]  
    word_sent= ' '.join([snow.stem(word) for word in word_sent])  
    word_sent = ' '.join([w for w in word_sent.split() if len(w)>1])  
    return word_sent
```

```
df_text['ProcessedText'] = df_text['Text'].astype(str).apply(preprocessing)
```

Data Preprocessing

Final Dataframe


	Positive	Text	ProcessedText
0	1	I have bought several of the Vitality canned d...	bought sever vital can dog food product found ...
1	0	Product arrived labeled as Jumbo Salted Peanut...	product arriv label jumbo salt peanut peanut a...
2	1	This is a confection that has been around a fe...	confect around centuri light pillowi citrus ge...
3	0	If you are looking for the secret ingredient i...	look secret ingredi robitussin believ found go...
4	1	Great taffy at a great price. There was a wid...	great taffi great price wide assort yummi taff...
...
393928	1	Great for sesame chicken..this is a good if no...	great sesam chicken good not better restur eat...
393929	0	I'm disappointed with the flavor. The chocolat...	disappoint flavor chocol note espec weak milk...
393930	1	These stars are small, so you can give 10-15 o...	star small give one train session tri train do...
393931	1	These are the BEST treats for training and rew...	best treat train reward dog good groom lower c...
393932	1	I am very satisfied ,product is as advertised,...	satisfi product advertis use cereal raw vinega...

393933 rows × 3 columns


Data Preprocessing

Finding the most frequent unigrams and bigrams

Top 10 Most Common Uni-grams



Top 10 Most Common bi-grams



generally positive

Data Preprocessing

Visualize the most frequent words using wordcloud



positive reviews (Score 4-5)



negative reviews (Score 1-3)

04

Models Creation





Models Creation

Stratified K-fold Cross Validation

```
from sklearn.model_selection import StratifiedKFold
skf = StratifiedKFold(n_splits=4, shuffle=True, random_state=1)
lst_accu_stratified = []
```

Create a Naive Bayes model

```
from sklearn import naive_bayes
Naive = naive_bayes.MultinomialNB()

for train_index, test_index in skf.split(X, y):
    X_train_fold, X_test_fold = X[train_index], X[test_index]
    y_train_fold, y_test_fold = y[train_index], y[test_index]

    Tfidf_vect = TfidfVectorizer(stop_words='english', max_df=0.8, dtype= np.float32)
    Tfidf_vect.fit(X_train_fold)

    X_train_Tfidf = Tfidf_vect.transform(X_train_fold)
    X_test_Tfidf = Tfidf_vect.transform(X_test_fold)
    Naive.fit(X_train_Tfidf, y_train_fold)
    y_pred = Naive.predict(X_test_Tfidf)
    lst_accu_stratified.append(f1_score(y_test_fold, y_pred))
```



Models Creation

Create a model of SGDClassifier

(Suitable for large and sparse dataset)

```
from sklearn.linear_model import SGDClassifier
sgd = SGDClassifier()

lst_accu_stratified_sgd = []
for train_index, test_index in skf.split(X, y):
    X_train_fold, X_test_fold = X[train_index], X[test_index]
    y_train_fold, y_test_fold = y[train_index], y[test_index]

    Tfidf_vect = TfidfVectorizer(stop_words='english', max_df=0.8, dtype= np.float32)
    Tfidf_vect.fit(X_train_fold)

    X_train_Tfidf = Tfidf_vect.transform(X_train_fold)
    X_test_Tfidf = Tfidf_vect.transform(X_test_fold)
    sgd.fit(X_train_Tfidf, y_train_fold)

    y_pred = sgd.predict(X_test_Tfidf)
    lst_accu_stratified_sgd.append(f1_score(y_test_fold, y_pred))
```



Models Creation

F1 Score of the Naive Bayes Model after Cross Validation

List of possible F1 score: [0.8914590332213512, 0.8911640381307757, 0.8911468261602047, 0.8909237153781592]

Maximum F1 score That can be obtained from this model is: 89.14590332213513 %

Minimum F1 score: 89.09237153781592 %

Overall F1 score: 89.11734032226228 %

Standard Deviation is: 0.00021963800599970014

**Small standard deviation
→ robust models**

F1 Score of the SGDClassifier Model after Cross Validation

List of possible F1 score: [0.918450211579385, 0.9183829511325337, 0.9182471484052476, 0.917640304642314]

Maximum F1 score That can be obtained from this model is: 91.8450211579385 %

Minimum F1 score: 91.76403046423141 %

Overall F1 score: 91.81801539398701 %

Standard Deviation is: 0.00036967704485381517





Models Creation

Dimensionality Reduction

Baseline SGDClassifier model:
73780 features ;
F1 Score: 91.81802%

TruncatedSVD transformer is often used on count/tf-idf matrices, known as latent semantic analysis (LSA).

```
from sklearn.decomposition import TruncatedSVD
svd = TruncatedSVD(n_components = 100) # n_components = 100, as recommended in sklearn documentation

X_train_svd = svd.fit_transform(X_train_Tfidf_svd, y_train)
X_test_svd = svd.transform(X_test_Tfidf_svd)
```

100 components:
F1 Score: 89.38094%

```
svd1 = TruncatedSVD(n_components = 500) # try to increase the number of components

X_train_svd1 = svd1.fit_transform(X_train_Tfidf_svd, y_train)
X_test_svd1 = svd1.transform(X_test_Tfidf_svd)
```

500 components:
F1 Score: 90.95921%



Findings

(Unigram TD-IDF)



	Features	Importance
1	great	3.987900
2	best	3.344341
3	love	3.066141
4	delici	3.026104
5	perfect	2.716438
6	excel	2.489259
7	amaz	2.154049
8	good	2.134706
9	awesom	1.851652
10	favorit	1.830844



	Features	Importance
1	disappoint	-4.488774
2	worst	-3.522368
3	ok	-3.468223
4	return	-3.306934
5	aw	-3.040941
6	unfortun	-2.956783
7	terribl	-2.927215
8	horribl	-2.865595
9	okay	-2.714790
10	bland	-2.639161

Findings

(Bigram TD-IDF)



Features		Importance
1	high recommend	0.690828
2	pleasant surpris	0.428914
3	tast great	0.350020
4	great product	0.344050
5	definit buy	0.335500
6	great tast	0.319084
7	love stuff	0.310013
8	bast tast	0.300389
9	realli good	0.299654
10	far best	0.292919



Features		Importance
1	wast money	-5.883018
2	wo buy	-2.909699
3	disappoint product	-1.438679
4	threw away	-1.405127
5	buyer bewar	-1.340235
6	bad batch	-1.300711
7	wo order	-1.261163
8	throw away	-1.220730
9	tast bad	-1.210144
10	tast ok	-1.192638

Models Creation



Hugging Face



LiYuan/**amazon-review-sentiment-analysis**



like

0



Text Classification



PyTorch

TensorBoard



Transformers

apache-2.0

bert

generated_from_train



Model card

Files and versions



Training metrics



Community



Edit model card

distilbert-base-uncased-finetuned-mnli-amazon-query-shopping

This model is a fine-tuned version of [nlptown/bert-base-multilingual-uncased-sentiment](#) on an [Amazon US Customer Reviews Dataset](#). The code for the fine-tuning process can be found [here](#). This model is uncased: it does not make a difference between english and English. It achieves the following results on the evaluation set:

Models Creation



```
from transformers import AutoTokenizer, AutoModelForSequenceClassification

tokenizer = AutoTokenizer.from_pretrained("LiYuan/amazon-review-sentiment-analysis")

model = AutoModelForSequenceClassification.from_pretrained("LiYuan/amazon-review-sentiment-analysis")
```

```
for index, entry in enumerate(corpus['text']):
    batch = tokenizer(entry, padding = True, truncation = True, max_length = 512, return_tensors = "pt")
    with torch.no_grad():
        outputs = model(**batch)
        predictions = F.softmax(outputs.logits, dim = 1 )
        labels = torch.argmax(predictions, dim = 1)
        labels = [model.config.id2label[label_id] for label_id in labels.tolist()]
        corpus.loc[index, 'labels'] = labels
```


Models Creation



label	label_hf	text
0	5	5 stars have bought several of the Vitality canned dog food products and have found them all to be of good quality. The product looks more like a stew than a processed meat and it smells better. My Labr...
1	1	1 star Product arrived labeled as Jumbo Salted Peanuts...the peanuts were actually small sized unsalted. Not sure if this was an error or if the vendor intended to represent the product as "Jumbo".
2	4	5 stars This is a confection that has been around a few centuries. It is a light, pillowy citrus gelatin with nuts – in this case Filberts. And it is cut into tiny squares and then liberally coated with ...
3	2	5 stars If you are looking for the secret ingredient in Robitussin I believe I have found it. I got this in addition to the Root Beer Extract I ordered (which was good) and made some cherry soda. The fl...
4	5	5 stars Great taffy at a great price. There was a wide assortment of yummy taffy. Delivery was very quick. If your a taffy lover, this is a deal.

Models Creation



```
mean_absolute_error(corpus['label'], corpus['label_hf'])
```

```
0.3682
```

```
f1_score(corpus['label'], corpus['label_hf'], average = None)
```

```
array([0.69375, 0.38857143, 0.41991925, 0.38023451, 0.88294314])
```

```
f1_score(corpus['label'], corpus['label_hf'], average = 'micro')
```

```
0.7444
```



06

Conclusion & Future Improvements

Conclusion & Future Improvements

Limitations

- Dataset is pretty large, requires long time to process
- No real-time tracking as Amazon.com cannot be scrapped directly



Future Improvements

- Apply bagging/ big data techniques
- Look for other data sources which updates regularly for real-time tracking





Next Steps

1

Better Visualization

Using BI tools e.g. Tableau to create a dashboard that make it more user-friendly

2

Compare with Competitors

Conduct similar sentiment analysis on competitors to know the opportunities and threats



THANKS!

